

# National Institute of Technology Karnataka Surathkal

## Department of Information Technology



### IT 301 Parallel Computing

Shared Memory Programming Technique (6)

OpenMP : *sections, threadprivate, collapse*

**Dr. Geetha V**

*Assistant Professor*

*Dept of Information Technology*

*NITK Surathkal*

# Index

- OpenMP

- Directives : if, for, master, single, Barrier, atomic, critical,
- Directives
  - Sections
- Clauses
  - Threadprivate
  - Collapse
  - Threadwait
  - Copyin, copyprivate

- References

# Course Outline

## Course Plan: Theory:

### Part A: Parallel Computer Architectures

Week 1,2,3: **Introduction to Parallel Computer Architecture:** Parallel Computing, Parallel architecture, bit level, instruction level , data level and task level parallelism. Instruction level parallelism: pipelining(Data and control instructions), scalar and superscalar processors, vector processors. Parallel computers and computation.

Week 4,5: Memory Models: UMA, NUMA and COMA. Flynn's classification, Cache coherence,

Week 6,7: Amdahl's Law. Performance evaluation, Designing parallel algorithms : Divide and conquer, Load balancing, Pipelining.

Week 8 -11: **Parallel Programming techniques like Task Parallelism using TBB, TL2, Cilk++ etc. and software transactional memory techniques.**

# Course Outline

## Part B: OpenMP/MPI/CUDA

- Week 1,2,3 : **Shared Memory Programing Techniques:** Introduction to OpenMP : Directives: *parallel, for, sections, task, master, single, critical, barrier, taskwait, atomic.* Clauses: *private, shared, firstprivate, lastprivate, reduction, nowait, ordered, schedule, collapse, num\_threads, if(), threadprivate, copyin, copyprivate*
- Week 4,5: **Distributed Memory programming Techniques:** MPI: Blocking, Non-blocking.
- Week 6,7 : CUDA : OpenCL, Execution models, GPU memory, GPU libraries.
- Week 10,11,: **Introduction to accelerator programming using CUDA/OpenCL and Xeon-phi. Concepts of Heterogeneous programming techniques.**

### Practical:


Implementation of parallel programs using OpenMP/MPI/CUDA.

**Assignment:** Performance evaluation of parallel algorithms (in group of 2 or 3 members)

# 1. OpenMP

## FORK – JOIN Parallelism

- OpenMP program begin as a single process: the master thread. The master thread executes sequentially until the first parallel region construct is encountered.
- When a parallel region is encountered, master thread
  - Create a group of threads by FORK.
  - Becomes the master of this group of threads and is assigned the thread id 0 within the group.
- The statement in the program that are enclosed by the parallel region construct are then executed in parallel among these threads.
- JOIN: When the threads complete executing the statement in the parallel region construct, they synchronize and terminate, leaving only the master thread.



## 2. OpenMP Programming: Worksharing

Work sharing constructs

- Loop constructs
- Section construct
- Single construct

## 2. OpenMP Programming: Section

```
#pragma omp sections [clause,...]
{
  [#pragma omp section new-line
    Structured block]
  [#pragma omp section new-line
    Structured-block]
```

### Clauses

Private(list)

Firstprivate(list)

Lastprivate(list)

Reduction(operator:list)

nowait

### Section :

- It is a non-iterative work-sharing construct that contains a set of structured blocks that are to be divided among, and executed by, the threads in a team.
- Each structured block is executed by one of the threads in the team.
- There is an implicit barrier at the end of sections construct, unless a *nowait* clause is specified.
- Only a single *nowait* clause can appear on a sections directive.

## 2. OpenMP Programming: Collapse

### Collapse :

```
#pragma omp for schedule(static, n)
collapse(2)
for(i=0; i<imax; i++) {
    for(j=0; j<jmax; j++)
        a[i][j] = b[i][j] + c[i][j]
}
```

- It increases the total number of iterations that will be partitioned across the available number of OMP threads by reducing the granularity of work to be done by each thread.
- If the amount of work to be done by each thread is non-trivial (after collapsing is applied), this may improve the parallel scalability of the OMP applications.



## 2. OpenMP Programming: threadprivate

```
#pragma omp threadprivate(list)
```

### Threadprivate

- Each thread is allowed to have its own temporary view of the shared memory.
- Each thread also has access to another type of memory that must not be accessed by other threads, called threadprivate memory
- **Shared variable:** each thread refers to the original variable.
- **Private variable:** Current thread's private version of the original variable.
- **Threadprivate:** variable appearing in threadprivate directives are threadprivate.

## 2. OpenMP Programming: threadprivate

### Threadprivate

```
#pragma omp threadprivate(list)
```

- It specifies that named global-lifetime objects are replicated, with each thread having its own copy.
- Each copy of the *threadprivate* object is initialized once.
- A thread may not reference another thread's copy of a *threadprivate* object.
- A *threadprivate* object must not appear in any clause except the *copyin*, *copyprivate*, *schedule*, *num\_threads*, and *if* clauses.
- The list is a comma-separated list of file-scope, names-scope, or static block-scope variables that do not have incomplete types.

## 2. OpenMP Programming: threadprivate

### Private vs Threadprivate

```
#pragma omp threadprivate(list)
```

- Private variable scope is defined for only specific parallel region. *Threadprivate* is variable scope is declared across the parallel regions.
- *Firstprivate()* is used to copy the values from original variable. *Copyin()* is used to copy the values while entering into parallel region first time.
- Private variables are stored on stack most of the time. *Threadprivate* variables are stored in heap or thread local storage.

## 2. OpenMP Programming: threadprivate

### Data Copying clauses

- These clauses support the copying of data values from private or *threadprivate* variables on one implicit task or thread to the corresponding variables on other implicit tasks or threads in the team.
- **Copyin(list)**: Copies the value of the master thread's *threadprivate* variable to the *threadprivate* variable of each other member of the team executing the parallel region.
- **Copyprivate (list)** : Broadcasts a value from the data environment of one implicit task to the data environments of the other implicit tasks belonging to the parallel region.

```
#pragma omp threadprivate(list)
```

```
Copyin(list)
```

```
Copyprivate(list)
```

## 2. OpenMP Programming: Examples- sections

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main (void) {
    int x=10, y=20, a=0, b=0;
    printf("1. x=%d,y=%d,a=%d,b=%d\n",x,y,a,b);
    #pragma omp parallel
    {
        #pragma omp sections
        {
            #pragma omp section
            {
                a=x+y;
                int tid1=omp_get_thread_num();
                printf("tid=%d,x+y=%d\n",tid1,a);
            }
            #pragma omp section
            {
                b=x*y;
                int tid2=omp_get_thread_num();
                printf("tid=%d, x*y=%d\n",omp_get_thread_num(),b);
            }
        }
    }
    printf("2.x=%d,y=%d,a=%d,b=%d\n",x,y,a,b);
    return 0;
}
```

```
1. x=10,y=20,a=0,b=0
tid=5,x+y=30
tid=2, x*y=200
2.x=10,y=20,a=30,b=200
```

Section 1 executes x+y

Section 2 executes x\*y

## 2. OpenMP Programming: Examples – threadprivate, copyin

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int count=0;
#pragma omp threadprivate(count)

int main (void) {
int x=10, y=20,a[10],b[10],c[10],i;
//int count=0;
for(i=0;i<10;i++)
    b[i]=c[i]=i;

printf("1. count=%d\n",count);
#pragma omp parallel num_threads(2) copyin(count)
{
    #pragma omp for schedule(static,5) firstprivate(x)
    for(i=0;i<10;i++)
    {
        int tid1=omp_get_thread_num();
        a[i]=b[i]+c[i];
        count++;
        x++;
        printf("tid=%d,a[%d]=%d, count=%d x=%d\n",tid1,i,a[i],count,x);
    }

    #pragma omp barrier
    printf("2. count=%d x=%d tid=%d\n",count,x,omp_get_thread_num());
}
```

```
#pragma omp for schedule(static,5) firstprivate(x)
for(i=0;i<10;i++)
{
    int tid1=omp_get_thread_num();
    a[i]=b[i]*c[i];
    count++;
    x++;
    printf("tid=%d,a[%d]=%d, count=%d, x=%d\n",tid1,i,a[i],count,x);
}

#pragma omp barrier
printf("4. count=%d x=%d\n",count,x);
printf("\n");
return 0;
}
```

Thread private and copyin  
- count value is retained  
across parallel regions  
Firstprivate similar  
to copyin(), but the  
difference is that the scope  
of the variable is private()  
to parallel region.

```
1. count=0
tid=0,a[0]=0, count=1 x=11
tid=0,a[1]=2, count=2 x=12
tid=1,a[5]=10, count=1 x=11
tid=1,a[6]=12, count=2 x=12
tid=1,a[7]=14, count=3 x=13
tid=1,a[8]=16, count=4 x=14
tid=1,a[9]=18, count=5 x=15
tid=0,a[2]=4, count=3 x=13
tid=0,a[3]=6, count=4 x=14
tid=0,a[4]=8, count=5 x=15
2. count=5 x=10 tid=0
tid=0,a[0]=0, count=6, x=11
tid=0,a[1]=1, count=7, x=12
2. count=5 x=10 tid=1
tid=1,a[5]=25, count=6, x=11
tid=1,a[6]=36, count=7, x=12
tid=1,a[7]=49, count=8, x=13
tid=1,a[8]=64, count=9, x=14
tid=1,a[9]=81, count=10, x=15
tid=0,a[2]=4, count=8, x=13
tid=0,a[3]=9, count=9, x=14
tid=0,a[4]=16, count=10, x=15
4. count=10 x=10
```



## 2. OpenMP Programming: Examples - Collapse

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main (void) {
    int i,j;
    #pragma omp parallel
    {
        #pragma omp for schedule(static,3) private(i,j)
        for(i=0;i<6;i++)
            for(j=0;j<5;j++)
            {
                int tid2=omp_get_thread_num();
                printf("tid=%d, i=%d j=%d\n",omp_get_thread_num(),i,j);
            }
    }

    return 0;
}
```

```
tid=0, i=0 j=0
tid=0, i=0 j=1
tid=0, i=0 j=2
tid=0, i=0 j=3
tid=0, i=0 j=4
tid=0, i=1 j=0
tid=0, i=1 j=1
tid=0, i=1 j=2
tid=0, i=1 j=3
tid=0, i=1 j=4
tid=0, i=2 j=0
tid=0, i=2 j=1
tid=0, i=2 j=2
tid=0, i=2 j=3
tid=0, i=2 j=4
tid=1, i=3 j=0
tid=1, i=3 j=1
tid=1, i=3 j=2
tid=1, i=3 j=3
tid=1, i=3 j=4
tid=1, i=4 j=0
tid=1, i=4 j=1
tid=1, i=4 j=2
tid=1, i=4 j=3
tid=1, i=4 j=4
tid=1, i=5 j=0
tid=1, i=5 j=1
tid=1, i=5 j=2
tid=1, i=5 j=3
tid=1, i=5 j=4
```

when the loop is executed  
without collapse ()

## 2. OpenMP Programming: Examples - Collapse

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>

int main (void) {
    int i,j;
    #pragma omp parallel
    {
        #pragma omp for schedule(static,3) private(i,j) collapse(2)
        for(i=0;i<6;i++)
            for(j=0;j<5;j++)
            {
                int tid2=omp_get_thread_num();
                printf("tid=%d, i=%d j=%d\n",omp_get_thread_num(),i,j);
            }
    }

    return 0;
}
```

When collapse is used both the for loops are used for scheduling.

```
tid=0, i=0 j=0
tid=0, i=0 j=1
tid=0, i=0 j=2
tid=0, i=4 j=4
tid=0, i=5 j=0
tid=0, i=5 j=1
tid=6, i=3 j=3
tid=6, i=3 j=4
tid=6, i=4 j=0
tid=3, i=1 j=4
tid=3, i=2 j=0
tid=3, i=2 j=1
tid=2, i=1 j=1
tid=2, i=1 j=2
tid=2, i=1 j=3
tid=7, i=4 j=1
tid=7, i=4 j=2
tid=7, i=4 j=3
tid=1, i=0 j=3
tid=1, i=0 j=4
tid=1, i=1 j=0
tid=1, i=5 j=2
tid=1, i=5 j=3
tid=1, i=5 j=4
tid=4, i=2 j=2
tid=4, i=2 j=3
tid=4, i=2 j=4
tid=5, i=3 j=0
tid=5, i=3 j=1
tid=5, i=3 j=2
```



## 2. OpenMP Programming: Examples – copyprivate()

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
int count=0;
#pragma omp threadprivate(count)

int main (void) {
int x=10, y=20, a[10], b[10], c[10], i;
//int count=0;
for(i=0; i<10; i++)
    b[i]=c[i]=i;

printf("1. count=%d\n", count);
#pragma omp parallel num_threads(2) copyin(count)
{
    #pragma omp for schedule(static,5) firstprivate(x)
    for(i=0; i<10; i++)
    {
        int tid1=omp_get_thread_num();
        a[i]=b[i]+c[i];
        count++;
        x++;
        printf("tid=%d, a[%d]=%d, count=%d x=%d\n", tid1, i, a[i], count, x);
    }

    #pragma omp barrier
    printf("2. before copyprivate count=%d x=%d tid=%d\n", count, x, omp_get_thread_num());
    #pragma omp single copyprivate(count)
    {
        count=count+20;
    }
    printf("3. after copyprivate count=%d x=%d tid=%d\n", count, x, omp_get_thread_num());

    #pragma omp for schedule(static,5) firstprivate(x)
    for(i=0; i<10; i++)
    {
        int tid1=omp_get_thread_num();
        a[i]=b[i]*c[i];
        count++;
        x++;
        printf("tid=%d, a[%d]=%d, count=%d, x=%d\n", tid1, i, a[i], count, x);
    }
}
#pragma omp barrier
printf("4. count=%d x=%d\n", count, x);
printf("\n");
return 0;
}
```

Copyprivate broadcast the value of the variable to all other threads

```
1. count=0
tid=0, a[0]=0, count=1 x=11
tid=0, a[1]=2, count=2 x=12
tid=0, a[2]=4, count=3 x=13
tid=0, a[3]=6, count=4 x=14
tid=0, a[4]=8, count=5 x=15
tid=1, a[5]=10, count=1 x=11
tid=1, a[6]=12, count=2 x=12
tid=1, a[7]=14, count=3 x=13
tid=1, a[8]=16, count=4 x=14
tid=1, a[9]=18, count=5 x=15
2. before copyprivate count=5 x=10 tid=1
2. before copyprivate count=5 x=10 tid=0
3. after copyprivate count=25 x=10 tid=1
tid=1, a[5]=25, count=26, x=11
tid=1, a[6]=36, count=27, x=12
tid=1, a[7]=49, count=28, x=13
tid=1, a[8]=64, count=29, x=14
tid=1, a[9]=81, count=30, x=15
3. after copyprivate count=25 x=10 tid=0
tid=0, a[0]=0, count=26, x=11
tid=0, a[1]=1, count=27, x=12
tid=0, a[2]=4, count=28, x=13
tid=0, a[3]=9, count=29, x=14
tid=0, a[4]=16, count=30, x=15
4. count=30 x=10
```

# Index

- OpenMP

- Directives : if, for, master, single, Barrier, atomic, critical,
- Directives
  - Sections
- Clauses
  - Threadprivate
  - Collapse
  - Threadwait
  - Copyin, copyprivate

- References

# Reference

## **Text Books and/or Reference Books:**

1. Professional CUDA C Programming – John Cheng, Max Grossman, Ty McKercher, 2014
2. B.Wilkinson, M.Allen, "Parallel Programming: Techniques and Applications Using Networked Workstations and Parallel Computers", Pearson Education, 1999
3. I.Foster, "Designing and building parallel programs", 2003
4. Parallel Programming in C using OpenMP and MPI – Micheal J Quinn, 2004
5. Introduction to Parallel Programming – Peter S Pacheco, Morgan Kaufmann Publishers, 2011
6. Advanced Computer Architectures: A design approach, Dezso Sima, Terence Fountain, Peter Kacsuk, 2002
7. Parallel Computer Architecture : A hardware/Software Approach, David E Culler, Jaswinder Pal Singh Anoop Gupta, 2011
8. Introduction to Parallel Computing, Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar, Pearson, 2011

# Reference

## Acknowledgements

1. Introduction to OpenMP <https://www3.nd.edu/~z xu2/acms60212-40212/Lec-12-OpenMP.pdf>
2. Introduction to parallel programming for shared memory Machines <https://www.youtube.com/watch?v=LL3TAHpxOig>
3. OpenMP Application Program Interface Version 2.5 May 2005
4. OpenMP Application Program Interface Version 5.0 November 2018

**Thank You**